



Search kaggle



Competitions

Datasets

Kernels

Discussion

Learn

**Gábor Vecsei**

Plant Seedlings Fun with Computer Vision

last run 4 months ago · Python notebook · 3689 views
using data from [Plant Seedlings Classification](#) · Public

54

voters



Tags

data visualization

data cleaning

image processing

plants

Notebook

Plant Seedlings Segmentation with pure Computer Vision

First of all, thanks for the popularity of this kernel. I hope it will help for you to create more accurate predictions

In [1]:

```
%matplotlib inline
import os
import matplotlib
import matplotlib.pyplot as plt
import pandas as pd
import cv2
import numpy as np
from glob import glob
import seaborn as sns
```

In [2]:

```
BASE_DATA_FOLDER = "../input"
TRAIN_DATA_FOLDER = os.path.join(BASE_DATA_FOLDER, "train")
```

Read images

First, I'll just read all the images. The images are in BGR (Blue/Green/Red) format because OpenCV uses this.

Btw... If you'd like to use RGB format, than you can use it, it won't effect the segmentation because we will use the HSV (Hue/Saturation/Value) color space for that.

```
In [3]:
images_per_class = {}
for class_folder_name in os.listdir(TRAIN_DATA_FOLDER):
    class_folder_path = os.path.join(TRAIN_DATA_FOLDER, class_folder_name)
    class_label = class_folder_name
    images_per_class[class_label] = []
    for image_path in glob(os.path.join(class_folder_path, "*.png")):
        image_bgr = cv2.imread(image_path, cv2.IMREAD_COLOR)
        images_per_class[class_label].append(image_bgr)
```

Number of images per class

```
In [4]:
for key,value in images_per_class.items():
    print("{0} -> {1}".format(key, len(value)))
```

```
Common wheat -> 221
Black-grass -> 263
Charlock -> 390
Scentless Mayweed -> 516
Maize -> 221
Cleavers -> 287
Loose Silky-bent -> 654
Small-flowered Cranesbill -> 496
Shepherds Purse -> 231
Sugar beet -> 385
Common Chickweed -> 611
Fat Hen -> 475
```


Plot images

Plot images so we can see what the input looks like

```
In [5]:
def plot_for_class(label):
    nb_rows = 3
    nb_cols = 3
    fig, axs = plt.subplots(nb_rows, nb_cols, figsize=(6, 6))


    n = 0
    for i in range(0, nb_rows):
        for j in range(0, nb_cols):
            axs[i, j].axis.set_ticklabels([])
            axs[i, j].yaxis.set_ticklabels([])
            axs[i, j].imshow(images_per_class[label][n])
            n += 1
```

```
In [6]:
plot_for_class("Small-flowered Cranesbill")
```



```
In [7]:
plot_for_class("Maize")
```





Preprocessing for the images:

Now comes the interesting and fun part!

I created separate functions so if you'd like to use these it is easier.

In the next block I'll explain what I am doing to make the segmentation happen.

In [8]:

```
def create_mask_for_plant(image):
    image_hsv = cv2.cvtColor(image, cv2.COLOR_BGR2HSV)
```

[Notebook](#) Code Data (1) Comments (19) Log Versions (6) Forks (11)

```
lower_hsv = np.array([00 + sensitivity, 100, 50],
                    dtype=np.uint8)
upper_hsv = np.array([60 + sensitivity, 255, 255],
                    dtype=np.uint8)

mask = cv2.inRange(image_hsv, lower_hsv, upper_hsv)
kernel = cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (11,11))
mask = cv2.morphologyEx(mask, cv2.MORPH_CLOSE, kernel)

return mask

def segment_plant(image):
    mask = create_mask_for_plant(image)
    output = cv2.bitwise_and(image, image, mask = mask)
    return output
```

Page 5 of 13

(<https://www.cs.auckland.ac.nz/courses/compsci773s1c/lectures/ImageProcessing-html/topic4.htm>).

Basically with the *Close* operation we would like to keep the shape of the original objects (1 blobs on the mask image) but close the small holes. That way we can clarify our detection mask more.




image from <https://www.cs.auckland.ac.nz/courses/compsci773s1c/lectures/ImageProcessing-html/topic4.htm>
(<https://www.cs.auckland.ac.nz/courses/compsci773s1c/lectures/ImageProcessing-html/topic4.htm>)

After these steps we created the mask for the object.

In [9]:


```
# Test image to see the changes
image = images_per_class["Small-flowered Cranesbill"][97]

image_mask = create_mask_for_plant(image)
image_segmented = segment_plant(image)
image_sharpen = sharpen_image(image_segmented)

fig, axs = plt.subplots(1, 4, figsize=(20, 20))
axs[0].imshow(image)
axs[1].imshow(image_mask)
axs[2].imshow(image_segmented)
axs[3].imshow(image_sharpen)
```

Out[9]:

```
<matplotlib.image.AxesImage at 0x7f204e735198>
```





After this step we can see that the image on the right is more recognizable than the original image on the left.

From the mask image what we created (because we need that for the segmentation), we can extract some features. For example we can see how the area of the plant changes based on their classes.

Of course from the contours we can extract much more information than the area of the contour and the number of components, but this is the one I would like to show you.

Additional read: https://en.wikipedia.org/wiki/Image_moment (https://en.wikipedia.org/wiki/Image_moment)

```
In [10]:
def find_contours(mask_image):
    return cv2.findContours(mask_image, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)[-2]

def calculate_largest_contour_area(contours):
    if len(contours) == 0:
        return 0
    c = max(contours, key=cv2.contourArea)
    return cv2.contourArea(c)

def calculate_contours_area(contours, min_contour_area = 250):
    area = 0
    for c in contours:
        c_area = cv2.contourArea(c)
        if c_area >= min_contour_area:
            area += c_area
    return area
```

```
In [11]:
areas = []
```

```

larges_contour_areas = []
labels = []
nb_of_contours = []
images_height = []
images_width = []

for class_label in images_per_class.keys():
    for image in images_per_class[class_label]:
        mask = create_mask_for_plant(image)
        contours = find_contours(mask)

        area = calculate_contours_area(contours)
        largest_area = calculate_largest_contour_area(contours)
        height, width, channels = image.shape

        images_height.append(height)
        images_width.append(width)
        areas.append(area)
        nb_of_contours.append(len(contours))
        larges_contour_areas.append(largest_area)
        labels.append(class_label)

```

In [12]:

```

features_df = pd.DataFrame()
features_df["label"] = labels
features_df["area"] = areas
features_df["largest_area"] = larges_contour_areas
features_df["number_of_components"] = nb_of_contours
features_df["height"] = images_height
features_df["width"] = images_width

```

In [13]:

```
features_df.groupby("label").describe()
```

Out[13]:

	area							
	count	mean	std	min	25%	50%	75%	max
label								
Black-grass	263.0	41793.982890	159636.952451	0.0	1559.750	7204.00	23556.250	2097133.5
Charlock	200.0	81660.004070	150000.000400	0.0	10705.105	20000.05	20700.105	1100000.5


Chenopodium	390.0	91669.094872	150626.968468	5177.5	12795.125	39382.25	80786.125	1193936.5
Cleavers	287.0	25619.073171	24726.375631	567.5	8368.000	18717.50	34363.000	175187.0
Common Chickweed	611.0	14935.878887	20172.494215	0.0	1917.500	5026.00	25140.000	131654.5
Common wheat	221.0	6949.357466	8161.181143	0.0	1883.000	4044.50	8953.000	57880.0
Fat Hen	475.0	30623.629474	67760.225738	512.0	2804.250	6153.50	28209.000	561161.5
Loose Silky-bent	654.0	24618.366208	132431.879535	0.0	718.500	1795.50	8252.500	1930040.5
Maize	221.0	115311.341629	206882.627747	851.5	3885.500	52033.00	119356.000	1435536.5
Scentless Mayweed	516.0	15517.381783	27678.754205	0.0	1589.250	3314.25	14044.125	198328.5
Shepherds Purse	231.0	38187.586580	66440.811037	562.0	3303.500	7435.50	45623.250	415662.5
Small-flowered Cranesbill	496.0	38126.337702	48735.741208	678.5	7834.625	16201.50	42556.500	306815.5
Sugar beet	385.0	79003.370130	112091.057669	0.0	16029.000	40748.50	81174.000	815699.0

12 rows × 40 columns

In [14]:

Did you find this Kernel useful?
Show your appreciation with an upvote

54




Comments (19)

All Comments

Sort by

Hotness



Click here to enter a comment...



leejet • Posted on Version 5 • 4 months ago • Options • Reply

^ 2 v



Actually, cv2.imread() will read the image in BGR mode. So, you read the image in BGR mode and then convert it to RGB mode. That's why you could show the normal image in pyplot which deal with image in RGB mode.



ABHISHEKDOG... • Posted on Version 3 • 5 months ago • Options • Reply

^ 0 v

This is a good step to highlight only the green areas for recognition. It will speed up the computation and be more accurate. Thanks



Gábor Vecsei • Posted on Version 4 • 4 months ago • Options • Reply

^ 1 v



Thanks! Btw, it does not speed up the process (because you have to convert the image from one color space to the other) but it does make your model more accurate because it doesn't have to deal with the background. By "more accurate" I mean a few decimals, but if you are in the first 10-20 in the contest this preprocessing can mean a lot.



ABHISHEKD... • Posted on Version 5 • 4 months ago • Options • Reply

^ 0 v

What if we grayscale or binarize the image after we get the green and black areas, I think it might speed up the computation as well?



Gábor Vecsei • Posted on Version 5 • 4 months ago • Options • Reply

^ 1 v



Yeah, it will, if you create or fine tune a model where the input shape accepts the grayscale images. But (for example) if you'd like to fine-tune Xception then the input shape has to be (299, 299, 3) and this means that it doesn't matter if the image is grayscale or not.



Jeru666 • Posted on Version 5 • 4 months ago • Options • Reply

^ 0 v

thanks for sharing this !!!



Carson S • Posted on Version 5 • 4 months ago • Options • Reply

^ 0 v

Nice post. Do you think you could give a little explanation of what you did and why for the preprocessing step?



Gábor Vecsei • Posted on Version 5 • 4 months ago • Options • Reply

^ 0 v

Of course, Later this day or tomorrow I'll write a more detailed explanation.



Gábor Vecsei • Posted on Latest Version • 4 months ago • Options • Reply

^ 0 v

I don't know if you noticed it, but I wrote a more detailed explanation.



Kalpana Krishna... • Posted on Latest Version • 4 months ago • Options • Reply

^ 0 v

Thanks for sharing, great place to get started



ArsenyAntonov • Posted on Latest Version • 4 months ago • Options • Reply

^ 0 v

Thanks for your kernel. Why did you decide you use sharpen filter here?



Gábor Vecsei • Posted on Latest Version • 4 months ago • Options • Reply

^ 2 v



Because as I saw, the images of the plants are blurry and I thought that that sharpening the segmented image would help bring up more details, so I can do the classification more accurately.

Btw. This comes from a previous experience with emotion detection from face images. I had 48x48 grayscale images there and when I applied a sharpening filter the accuracy increased by a few points for the same model.



Siddique • Posted on Latest Version • 2 months ago • Options • Reply

^ 0 v

Many thanks for the code. It was very nice to see the background noise gone. Improved classification accuracy by 5+%.



Gábor Vecsei • Posted on Latest Version • 2 months ago • Options • Reply

^ 0 v

It is really nice to hear that it helped you. If you've got any questions, don't hesitate to ask!



Jacken • Posted on Latest Version • 12 days ago • Options • Reply

^ 0 v

I suppose how you can improve so a lot. Can you give me some advice?



Mohammed Alb... • Posted on Latest Version • 2 months ago • Options • Reply

^ 0 v

Excellent work and great details. It helped me a lot, many thanks.

I trained two models, one with and the other without the background filters and it improved the results a lot from 0.78 to 0.81 on the same CNN layer. Still trying to get a better score.



城璐 • Posted on Latest Version • a month ago • Options • Reply

^ 0 v

Many thanks for the tricks!



AdityaMishra • Posted on Latest Version • a month ago • Options • Reply

^ 0 v

Thanks!! The trick to sharpen the images & removing the background noise was really great!! Do point me to other kernels, if they use similar cool features for preprocessing .



Jacken • Posted on Latest Version • 12 days ago • Options • Reply



Good job! And I have used this trick by means of deeplabV3+ to get a good segmentation. However, I did not get an obvious improvement as others discussed these. I wonder what tricks you used, like param's value, etc. And I also want to know is this a good idea to other transfer learning task or fine tuning task? Thanks a lot.